

GitHub Recommender System

Big Data Computing Project
2020/2021

Fabrizio Rossi 1815023
Matteo Orsini 1795119




Goal

- **GitHub** is the largest source code host
 - 40+ million users
 - 200+ million repositories
- GitHub users can **star** interesting repositories
 - Used to suggest new useful repositories
- Our project goal is:
 - Given a list of starred repositories or an already existing user
 - Predict new repositories which may be interesting



Challenges

- Find suitable **dataset**:
 - We couldn't, so we collected it ourselves
- We don't have ratings: **implicit feedback**
- Implementation of models with PySpark
- Deal with a large amount of data




Data collection



Dataset collection

- Custom Python scripts with *requests* and *BeautifulSoup*
- Started from users that starred the *React.js* repository
- Collected repositories starred by them
- Collected repositories metadata
 - creator, name
 - about (short description)
 - main programming language
 - stars, forks
 - updated
 - sponsor (a form of donations)
- Total: 10.000 users and 354.981 repositories



Dataset Pruning




Dataset pruning I

- Some models require to process DataFrames with **100+ billions of records**
 - Even if number of users/repos is restricted, number of comparisons is huge
- Unfortunately neither Colab or Databricks can provide us the necessary computational power for free
- So, we **pruned** the dataset obtaining:
 - Users: 1.000
 - Repositories: 1.000
 - Starred relationships: 104.499



Dataset pruning II

- **How** we did it:
 - Took 1.000 users with highest number of starred repos
 - Took 1.000 most starred repos among those starred by the previous users
 - Discarded all user-repo relationships not relative to the previous users/repos
- Even with these numbers, still have many records:
 - Example: Item-Based Collaborative Filtering
 - Dataframe with `users x ns_repos x s_repos` records
 - Total: 64.413.366



Models



Content-Based Filtering I

- Based on **features** extracted from repositories
 - Doesn't need to know ratings from other users
- Needs **feature engineering**:
 - Text processing for the *about* field
 - lowercase, punctuation, tokenization
 - stopwords
 - stemming
 - tf-idf
 - One-Hot Encoding for the *language* field
 - Time conversion for the *updated* field
 - MinMax Normalization for the fields *stars*, *forks*, *updated*
- Build **item profiles**



Content-Based Filtering II

- Build **user profiles**
 - Prepare for each user the vectors of the starred repositories
 - Compute the user profile as the average of those vectors
 - We used the Summarizer class offered by PySpark
 - *groupBy* user id and compute mean of starred repos vectors
 - It works even with sparse vectors, so we can save some RAM
- Compute **cosine similarity** between each user profile and all item profiles of the repos that the user has not rated



Content-Based Filtering III

- Take top-k most similar repos for each user
 - We used a Window function
 - Partition by user id
 - Sort by similarity
 - Filter by row number

user_id	repo_id	similarity
1	2	0.92
1	3	0.56
1	4	0.88
2	1	0.92
2	3	0.31
2	4	0.25



Content-Based Filtering III

- Take top-k most similar repos for each user
 - We used a Window function
 - Partition by user id
 - Sort by similarity
 - Filter by row number

user_id	repo_id	similarity
1	2	0.92
1	3	0.56
1	4	0.88
2	1	0.92
2	3	0.25
2	4	0.31



Content-Based Filtering III

- Take top-k most similar repos for each user
 - We used a Window function
 - Partition by user id
 - Sort by similarity
 - Filter by row number

user_id	repo_id	similarity
1	2	0.92
1	4	0.88
1	3	0.56
2	1	0.92
2	4	0.31
2	3	0.25



Content-Based Filtering III

- Take top-k most similar repos for each user
 - We used a Window function
 - Partition by user id
 - Sort by similarity
 - Filter by row number

user_id	repo_id	similarity	row_number
1	2	0.92	1
1	4	0.88	2
1	3	0.56	3
2	1	0.92	1
2	4	0.31	2
2	3	0.25	3



Content-Based Filtering III

- Take top-k most similar repos for each user
 - We used a Window function
 - Partition by user id
 - Sort by similarity
 - Filter by row number

user_id	repo_id	similarity	row_number
1	2	0.92	1
1	4	0.88	2
1	3	0.56	3
2	1	0.92	1
2	4	0.31	2
2	3	0.25	3



User-Based Collaborative Filtering I

- No need to extract features from repositories
- Based on **user rating vectors**
 - Set of repositories starred by the user
- Steps:
 - Build user rating vectors
 - For each user, compute similarity among other user rating vectors
 - This time we used the **Jaccard similarity**
 - Take top-k most similar users (*Window function*)
 - Take repos starred by those users
 - Recommend them by assigning to each one the average rating given by other users



User-Based Collaborative Filtering II

- The **average rating** in our case is just the normalized count of top-k users that starred that repo:

$$r_{v,i} \in \{0, 1\} \qquad r_{u,i} = \frac{1}{k} \sum_{v \in \mathcal{U}^k} r_{v,i}$$



Item-Based Collaborative Filtering I

- Based on **item rating vectors**
 - Users that starred the repository
- Achieves better performances than User-Based but **high RAM usage**
- Steps:
 - Build item rating vectors
 - For each user we need to compute a score for each not starred repo:
 - Compute similarity with actually starred repos (**Jaccard similarity**)
 - Take top-k most similar starred repos (*Window function*)
 - Compute a score by aggregating the ratings of the found repos
 - Recommend non starred repos with highest obtained score



Item-Based Collaborative Filtering II


- In this case the **aggregation can't be performed** as a simple average since we have that all the ratings are 1
- We opted to use the **similarity between items** to compute the score, as follows:

$$r_{u,i} = \frac{\sum_{i' \in \mathcal{I}_u^k} \text{sim}(i, i')}{k}$$



Item-Based Collaborative Filtering III

- We have to **compute many times** the similarity between the same pairs of items
- So, we pre-computed the *Jaccard similarity* between each pair of item to make the computation faster



Matrix Factorization I

- **Ratings matrix** R seen as product of a user matrix X and an item matrix W :

$$R = X \times W^T$$

- X and W computed based on **hidden factors** extracted from observed starring relationships
- Score for a non-starred repo computed as dot product between user and repo vector

$$r_{u,i} = x_u^T \cdot w_i = \sum_{j=0}^d x_{u,j} w_{j,i}$$



Matrix Factorization II

- There are various methods to compute the two matrices
- **Alternating Least Squares (ALS)** method implemented and provided by PySpark
- Important parameters:
 - **rank**: rank of the latent matrices (5)
 - **implicitPrefs**: adapt implementation to implicit feedbacks (True)
 - **alpha**: feedback confidence (1.0)
- **Challenge**: how to predict repos for new user not in training?
 - Solution: non-standard approaches, not implemented in PySpark
 - So, in our case, train again the model (doesn't take much time)



Evaluation



Evaluation

- We computed the MAP@K and the Personalization measure
- **Mean Average Precision at K** (MAP@K) is the mean of the Average Precision at K (AP@K) metric computed for each user
 - Takes into consideration precision of the system and order of the items recommended
- **Personalization** is defined as the average of the dissimilarity between users lists of recommendations
 - Tells us if the recommender system produces items which are personalized for each user, or always the same items to different users



Evaluation II

Model	MAP@1	MAP@2	MAP@3	MAP@4	MAP@5	Personalization
Content-based	0.067	0.051	0.043	0.037	0.033	0.676
User-based	0.357	0.261	0.213	0.183	0.164	0.965
Item-based	0.388	0.308	0.263	0.229	0.205	0.680
Matrix Factorization	0.506	0.395	0.332	0.298	0.268	0.864

- **Content-Based Filtering** low performances probably due to low discriminance of extracted repos features
 - about field very short and contains recurring terms → similar TF-IDF vectors for different repos



Demo



Demo

- Built in **React** (frontend) + **Flask** (backend)
- Uses **ngrok** to host the server directly on Colab
- Predict repos starting from:
 - Already existing user of the system
 - A list of repos starred by a new user
- Available at <https://recommend-hub.netlify.app/>

